Same stream, drawn like an agent

What the Playground shows between send and answer, redrawn from the events it already receives. No new SSE event, no new field, and no backend work for the first step.

LOG ABOVE THE ANSWER	LINES ON SCREEN	STEP INSTRUCTION PRINTED	SITES VISIBLE WHILE WORKING	NEW OR CHANGED SSE EVENTS
~~1,640 px~~ **0 px**	~~17 rows~~ **6 lines**	~~9 times~~ **0 times**	~~0~~ **10 sites**	**0**
219 px at most while the agent works, 243 px with step 1 alone	At the busiest moment of the run	One label of a few words per step	When the calls return: 4 chips and "+6" in the feed, all 10 in Sources	Every state reads a field already on the wire

TODAY 0:25 into the recorded stock analysis run

The log continues above

Successfully Executed

Fetch the latest stock news (earnings reports, major announcements) for AAPL, TSLA, and AMZN; get the current price and recent historical price time series for each (for trends and support/resistance levels); and get RSI analysis and moving average analysis for each of the three stocks.

Agents execution failed

Failed

Fetch the latest stock news (earnings reports, major announcements) for AAPL, TSLA, and AMZN; get the current price and recent historical price time series for each (for trends and support/resistance levels); and get RSI analysis and moving average analysis for each of the three stocks.

Analyzing the prompt...

Using the gathered data, compose a report for AAPL, TSLA, and AMZN. For each stock include: News Summary, Technical Analysis (current price trend, key support/resistance levels, RSI and moving averages), and a clear Buy/Hold/Sell recommendation with a brief explanation. Send the report by email to you@company.com with an appropriate subject line.

report now, noting that AMZN news was rate-limited so I'll flag it as unavailable. For AAPL, I'm working out the latest price against the monthly range, RSI, and SMA9 to settle on a Hold recommendation with sup

PROPOSED The same run at the same moment, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 25s

Suggested plugins

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

Gathering news, prices and indicators 14 of 15 calls · 4.4 s

barchart.com · finance.yahoo.com · simplywall.st · fastcompany.com · +6

Emailing the report

SUMMARY AND DECISION

Ship the run view as a frontend change first. The prompt changes follow

Step 1 carries most of the visible change, needs no backend work and switches off with a flag. Steps 2 and 3 are independent of each other.

EXECUTIVE SUMMARY

- Today the Playground prints every status frame as a row. One recorded run drew 17 rows, 2 thinking cards and the step instruction 9 times: 1,640 px of log above the answer
- The stream already carries what an agent-style view needs: plugin logos, the plan sentence, each plugin call as it is written, page links and the email draft
- The proposed view draws the same frames as at most 6 short lines, folds them into the header when the answer starts, and fills the Outputs card and the Agent panel live
- Step 1 is frontend only. Step 2 adds three text-only prompt addenda for the wording. Step 3 is an optional wrapper change that brings every search favicon forward
- The terminal and to-do view for agent runs, the composer and the top bar do not change

THE BRIEF, POINT BY POINT

ASKED FOR	ANSWER	PAGE
No new events	0 events and 0 fields added	21
Feel like an agent before the terminal	One live line, a sentence, short step lines	6 to 14
No internal thinking, no plugin query	One word, “Thinking”. Labels instead of instructions	7, 12, 16
Short chunks saying what it is doing	Step labels and a call-by-call ticker	8, 15
A favicon for every site Perplexity reads	Favicon chips and a Sources list	16, 17
“Suggested plugins” with logos only	One label and a row of logos	6
The right side fills: Outputs and Agent	Outputs card and a run section in the Agent panel	11, 20
Subtle changes only	Same type, colours and layout. Behind a flag	19, 24

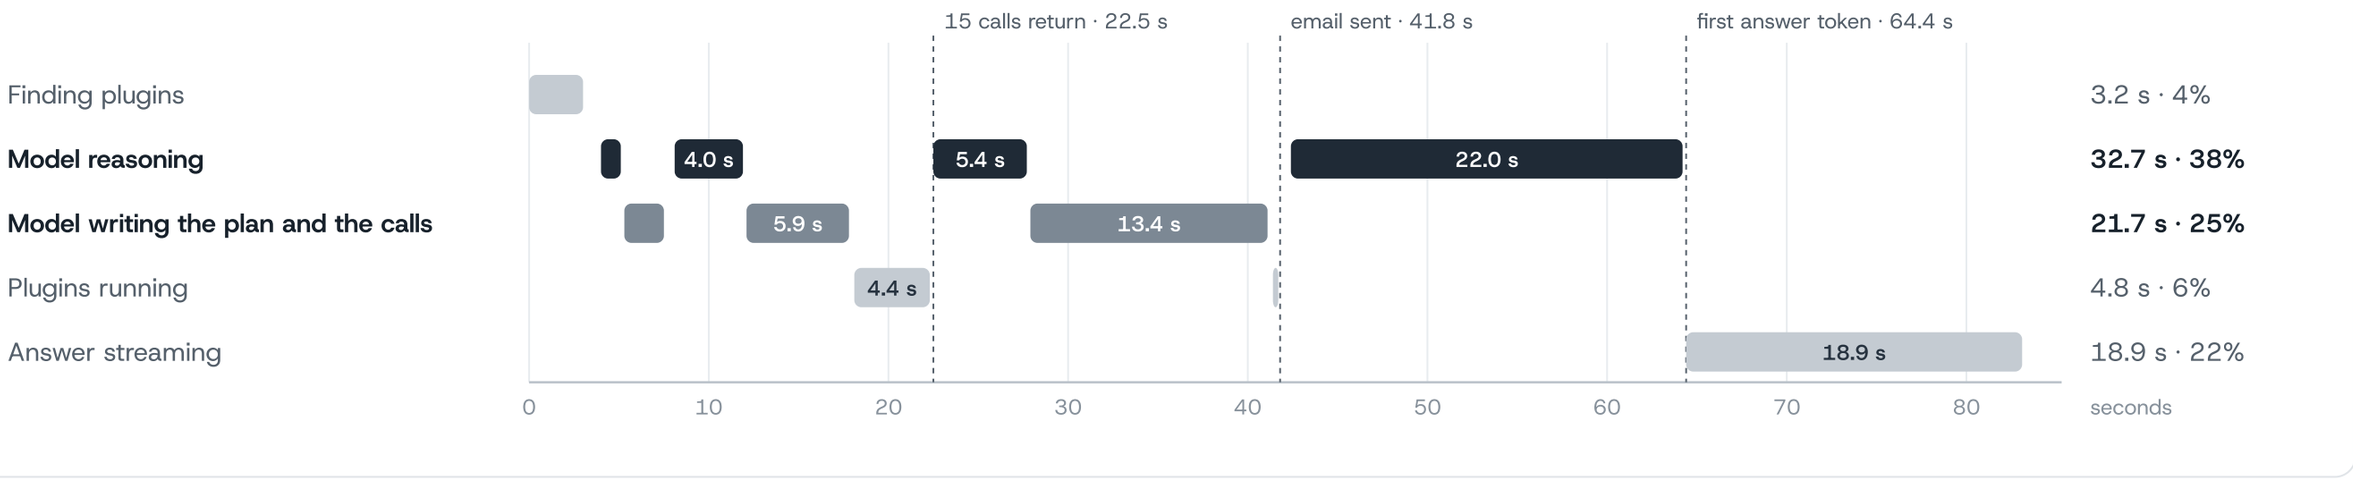
DECISION TABLE

STEP	WHAT SHIPS	OWNER	IMPACT 40%	EFFORT 25%	RISK 20%	REVERSIBLE 15%	WEIGHTED
1 · UI only	One reducer, the run feed, the Outputs card and the Agent panel section, behind a flag	Frontend	5	3	4	5	**4.30**
2 · Prompt addenda	Three short text blocks in the planner, step and answer prompts	Prompt owner	3	5	3	5	3.80
3 · Sources-first wrapper	The search wrapper returns links and titles before page text	Plugin owner	2	4	4	4	3.20

Scores are estimates. 5 is best in every column, so a high Effort or Risk score means little effort or little risk.

The wait is model time: 54 of the 64 seconds before the answer, against 5 for plugins

Stock analysis starter, recorded frame by frame on 5 Oct 2026: 1,033 frames in 85.3 s. Today the model's time shows as raw reasoning or as nothing at all. That is the time the new view makes readable.



Each bar is a stretch of the stream. A reasoning bar runs until the model starts writing. The remaining 4 s of the 85.3 s run are gaps between frames and the closing frames.

STRETCH	SECONDS	SHARE	TODAY THE SCREEN SHOWS	PROPOSED
Finding plugins	3.2	4%	Two status rows, then chips with full names	“Finding plugins”, then one row of logos
Model reasoning	32.7	38%	Raw reasoning text, memory status lines included	One word, “Thinking”, or the step label. The text is one click away
Model writing the plan and the calls	21.7	25%	Nothing new. The plan and the calls stream unseen	The opening sentence types in, chips appear call by call, the email draft fills Outputs
Plugins running	4.8	6%	Groups that repeat the logos and the instruction	Chips pulse: “15 calls running”. Then favicons
Answer streaming	18.9	22%	The answer, under 1,640 px of log	The answer, directly under the header

FRAMES

- 385 of 1,033 frames arrive on the `thinking` channel
- 242 of those are `step_output` : the plugin calls being written. Today they are not drawn

PLUGINS

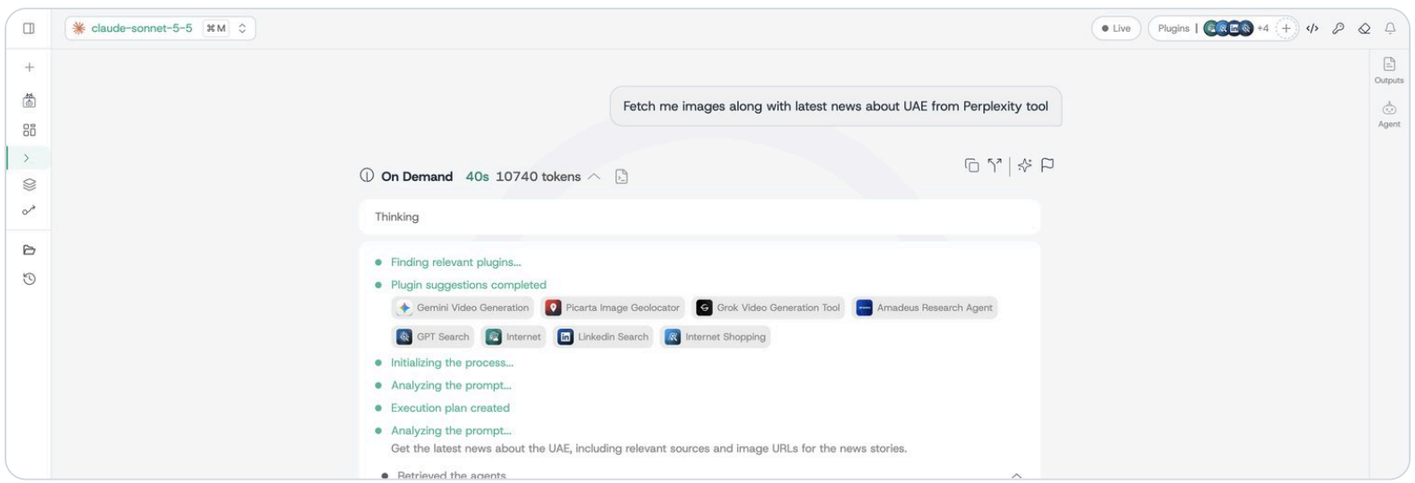
- 16 calls. 15 ran in parallel and returned together after 4.4 s
- 1 failed with HTTP 429, a rate limit at the data provider

SECOND RECORDING

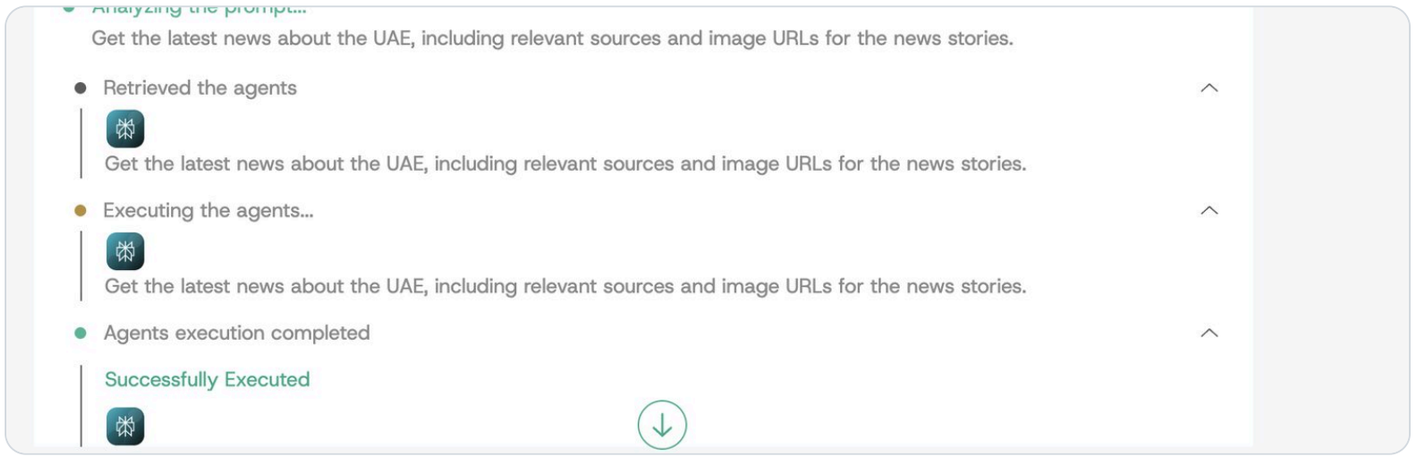
- Web search through Perplexity: 414 frames in 40.4 s
- One search call of 8.8 s. The answer cited 9 links on 7 sites

Today every status frame becomes a row, and the step instruction repeats under each one

Two of the screenshots supplied on 4 Oct 2026, cropped to the app. The “Today” panes on the following pages are a replica drawn from the recorded frames with the app's own type, colours and spacing.



Top of a run: a Thinking card, a row per status and suggestion chips with full names.



Further down the same run: the same step instruction under three consecutive rows.

MEASURED IN THE LIVE APP, STOCK ANALYSIS RUN

Status rows	17
Thinking cards with raw reasoning	2
Copies of the step instruction, about 50 words each	9
Log above the answer, on a 962 px screen	1,640 px
Suggestion chips with full names	5 to 8
Sites shown while working	0
Items in the Outputs panel during the run	0

WHAT THAT COSTS

- The answer starts more than a screen and a half below the header, and the log stays open after the run
- The user reads planner instructions and raw reasoning, which were written for the model
- For 21.7 s the model is writing the plan or the calls and nothing on screen changes
- The right side stays empty, and the Agent button opens configuration only

WHAT ALREADY WORKS AND STAYS

- The terminal and to-do view for agent runs
- The header with time and tokens, the Logs and Plan buttons
- Plugin logos, the composer and the top bar

REFERENCE · CHATGPT AND CLAUDE, CHECKED LIVE ON 5 OCT 2026

ChatGPT and Claude share one playbook: seven patterns, and the current stream supports all seven

Both were checked live while running multi-step tasks, and each pattern was seen in at least one of them. The proposal adopts the conventions and keeps what is OnDemand's own: plugin logos, the terminal view and the Outputs and Agent rail.

PATTERN	SEEN LIVE	ONDEMAND TODAY	PROPOSED
One live line with elapsed time	ChatGPT: “Working for 1m 43s”. Claude: “Batch — 4 actions · 1m 23s”	A header and a log that keeps growing	“On Demand · Working for 24s”
The agent says what it will do	ChatGPT opens with one first-person sentence before any step	Nothing. `objective` exists but is not drawn	`objective` as the opening sentence
Steps are short labels	ChatGPT: “Searching the web”. Claude: “Created 4 files, read 5 files, and 16 more steps”	The status message plus the full step instruction	A label of a few words per step, with plugin logo chips
Sites show as favicon chips	ChatGPT: a chip per site, then “Searched 24 websites”. Answer links carry a favicon	None	A favicon chip per returned page and a favicon on each answer link
Reasoning is one word until asked	ChatGPT: “Thinking” with a chevron	Raw reasoning streamed into the feed	“Thinking”, with the text one click away
The run folds when the answer arrives	ChatGPT: “Worked for 2m 16s”, which reopens the steps	The whole log stays open above the answer	The header keeps time and tokens. A chevron reopens the steps
A card on the right fills during the run	ChatGPT: Outputs and Sources. Claude: Progress, the tool in use, “Used in this session”	A docked Outputs panel that stays empty in plugin runs	An Outputs card with outputs and sources. A run section in the Agent panel

SIX RULES THE DESIGN FOLLOWS

Say what, not how

Labels of a few words. No instruction text and no state names such as “agents” or “fulfilling”

One line per step

A step updates in place and shrinks to one line when it is done

Logos and favicons carry the detail

Names show on hover and to screen readers

Reasoning is available, never pushed

One word while it runs, “Thought for 33s” afterwards

The answer owns the screen

The steps fold away when the first answer token lands

Nothing is lost

Plan, calls, failures and memory move to the right rail. The raw log stays behind Logs

Plugin suggestions shrink to a label and a row of logos

Same five plugins, same frame. The names move to hover, and the two status rows above them go.

TODAY Every status frame becomes a row

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand

- Finding relevant plugins…
- Plugin suggestions completed

Alchemy Agent · Chat Analytics Tool · Sentiment Analysis Report Tool · Quarterly financial report data tool · US Mutual Fund Technical

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 4s

Suggested plugins

Thinking

PROPOSED Outputs card

Outputs · Agent

OUTPUTS

Files, drafts and media appear here while the agent works

SOURCES

Pages the agent reads appear here, site by site

TODAY DRAWS

- Two status rows: “Finding relevant plugins…” and “Plugin suggestions completed”
- One chip per plugin with logo and full name, on two lines

PROPOSED DRAWS

- One line: “Suggested plugins” and five logos
- The plugin name shows on hover and to screen readers
- While the lookup runs, the live line reads “Finding plugins”

READ FROM THE STREAM

- `plugin_suggestion.initialized` at 0.0 s
- `plugin_suggestion.completed` at 3.2 s
- `suggestedPlugins[]` : `name` , `logoUrl`

The plan's own sentence opens the run while the plan is still being written

Today the plan streams unseen and the screen shows status rows and a Thinking card. The proposed view reads the objective out of that stream and types it in about two seconds before the plan frame lands.

TODAY Every status frame becomes a row

The log continues above

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand

Thinking

- Finding relevant plugins...
- Plugin suggestions completed
  - Alchemy Agent
  - Chat Analytics Tool
  - Sentiment Analysis Report Tool
  - Quarterly financial report data tool
  - US Mutual Fund Technical
- Initializing the process...
- Analyzing the prompt...

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 7s

Suggested plugins

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

PROPOSED Agent panel

Outputs | Agent

THIS RUN

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

MEMORY

0 memories used
Checked 1 time

CONFIGURATION

- AGENT
- MODEL
- LIVE
- SAMPLING
- SKILLS
- ADVANCED

TODAY DRAWS

- Rows: “Initializing the process...” and “Analyzing the prompt...”
- A Thinking card above the log. In this run it held only memory status lines
- The plan text itself is not drawn

PROPOSED DRAWS

- One sentence in plain words, typed as it streams
- No row for initializing or analyzing
- Reasoning stays behind the word “Thinking”

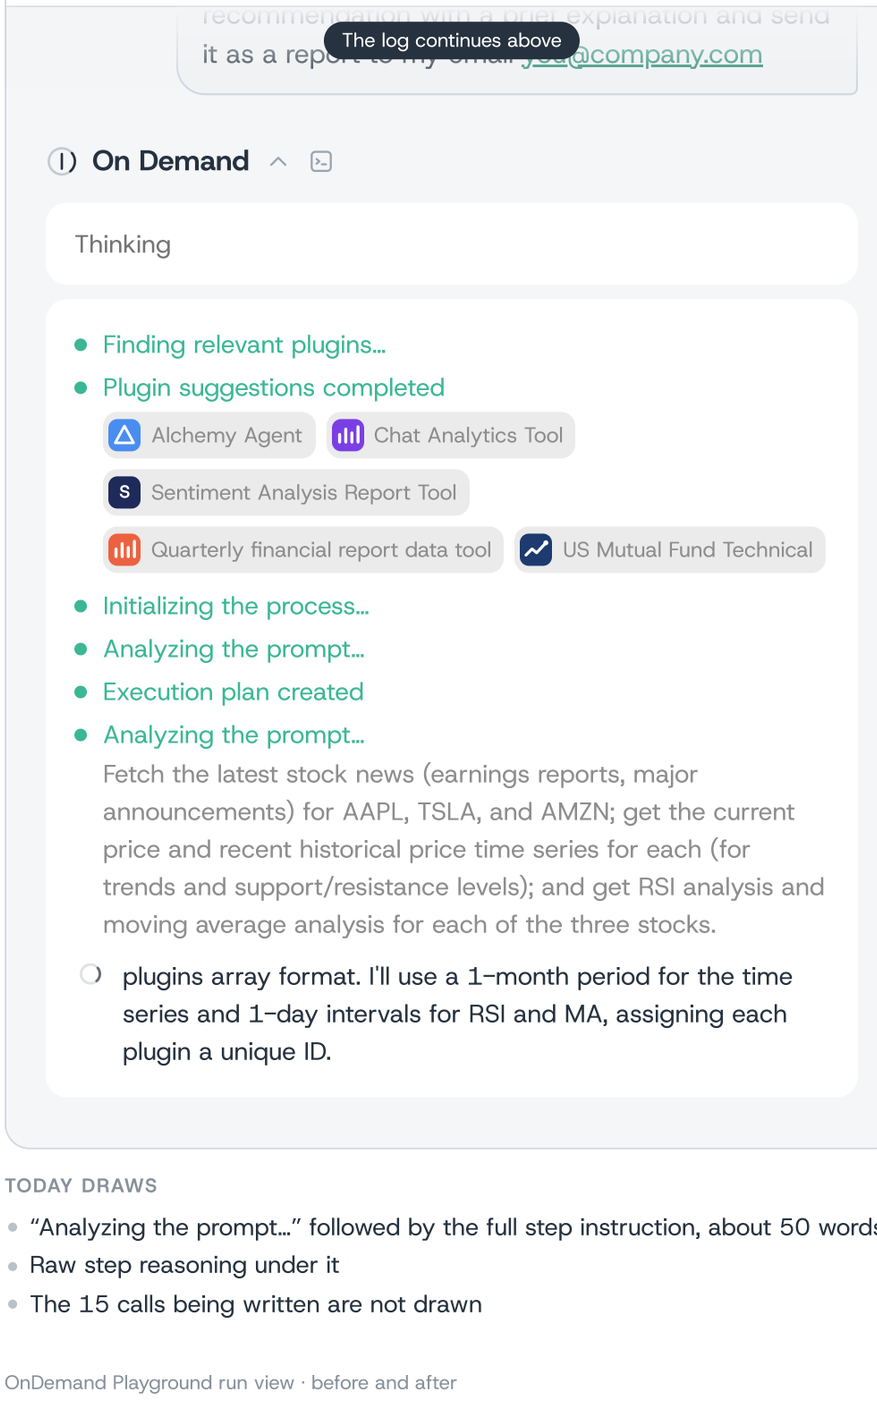
READ FROM THE STREAM

- `planning_thinking` 4.0 to 4.4 s
- `planning_output` 5.3 to 7.7 s: the plan JSON as text
- `plan_created` at 7.7 s confirms it

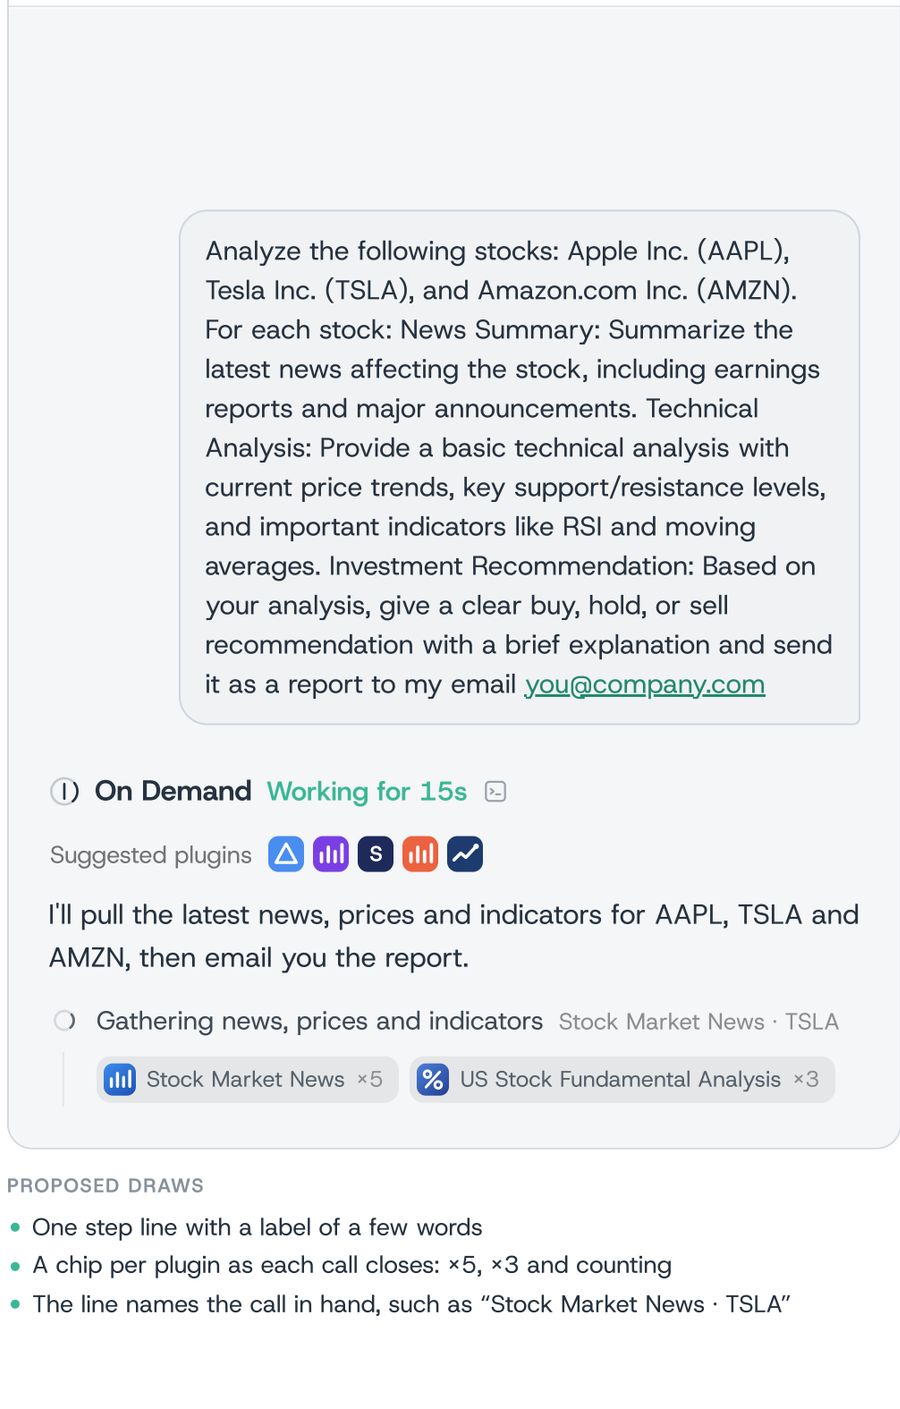
Plugin calls show up as logo chips while the model is still writing them

This stretch lasts 5.9 s. Today it shows the full step instruction and raw reasoning. Proposed: a short label, one chip per plugin with a counter that grows call by call, and the call in hand named beside the label.

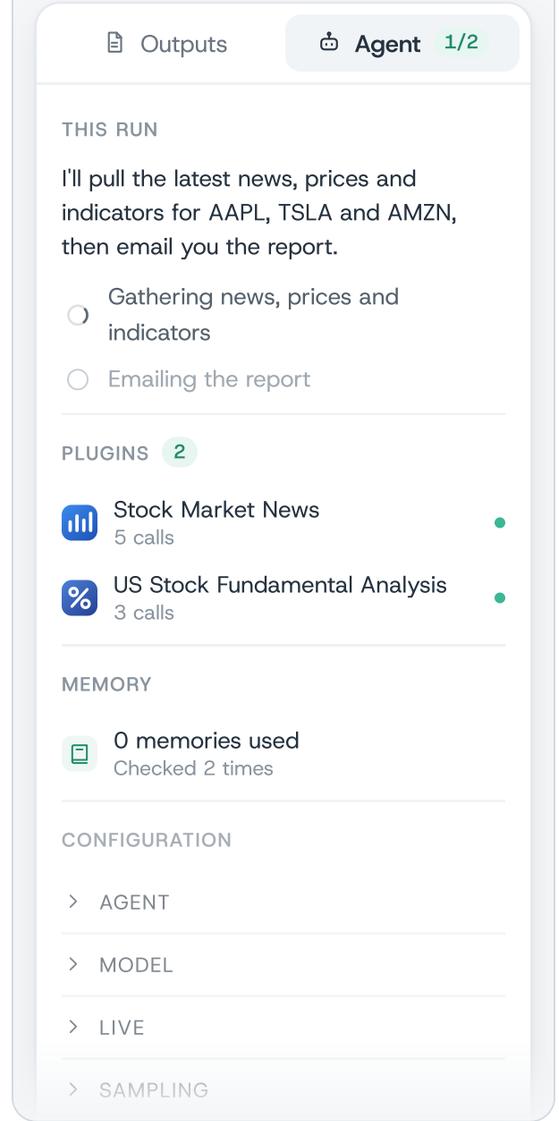
TODAY Every status frame becomes a row



PROPOSED The same run, with steps 1 and 2



PROPOSED Agent panel



TODAY DRAWS

- "Analyzing the prompt..." followed by the full step instruction, about 50 words
- Raw step reasoning under it
- The 15 calls being written are not drawn

PROPOSED DRAWS

- One step line with a label of a few words
- A chip per plugin as each call closes: ×5, ×3 and counting
- The line names the call in hand, such as "Stock Market News · TSLA"

READ FROM THE STREAM

- `analyzing` with `stepQuery` at 8.1 s
- `step_thinking` 8.1 to 12.1 s
- `step_output` 12.1 to 18.0 s: the calls JSON as text

Fifteen parallel calls read as one line instead of two groups

Today “Retrieved the agents” and “Executing the agents” each print 15 logos and the instruction again. Proposed: the chips pulse and the line reads “15 calls running”.

TODAY Every status frame becomes a row

The log continues above

Quarterly financial report data tool · US Mutual Fund Technical

- Initializing the process...
- Analyzing the prompt...
- Execution plan created
- Analyzing the prompt...

Fetch the latest stock news (earnings reports, major announcements) for AAPL, TSLA, and AMZN; get the current price and recent historical price time series for each (for trends and support/resistance levels); and get RSI analysis and moving average analysis for each of the three stocks.

- Retrieved the agents

Fetch the latest stock news (earnings reports, major announcements) for AAPL, TSLA, and AMZN; get the current price and recent historical price time series for each (for trends and support/resistance levels); and get RSI analysis and moving average analysis for each of the three stocks.

- Executing the agents...

Fetch the latest stock news (earnings reports, major announcements) for AAPL, TSLA, and AMZN; get the current price and recent historical price time series for each (for trends and support/resistance levels); and get RSI analysis and moving average analysis for each of the three stocks.

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 20s

Suggested plugins

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

Gathering news, prices and indicators 15 calls running

- Stock Market News ×6
- US Stock Fundamental Analysis ×3
- US Stock Technical Analysis ×6

PROPOSED Agent panel

Outputs · Agent 1/2

THIS RUN

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

- Gathering news, prices and indicators
- Emailing the report

PLUGINS 3

- Stock Market News — 6 calls
- US Stock Fundamental Analysis — 3 calls
- US Stock Technical Analysis — 6 calls

MEMORY

0 memories used
Checked 2 times

CONFIGURATION

AGENT
MODEL
LIVE

TODAY DRAWS

- Group “Retrieved the agents”: 15 logos and the instruction
- Group “Executing the agents...”: the same logos and the instruction again

PROPOSED DRAWS

- No new line. The draft chips become the confirmed call list
- Chips pulse while the calls run
- The Agent button shows 1/2 and the panel lists each plugin with its calls

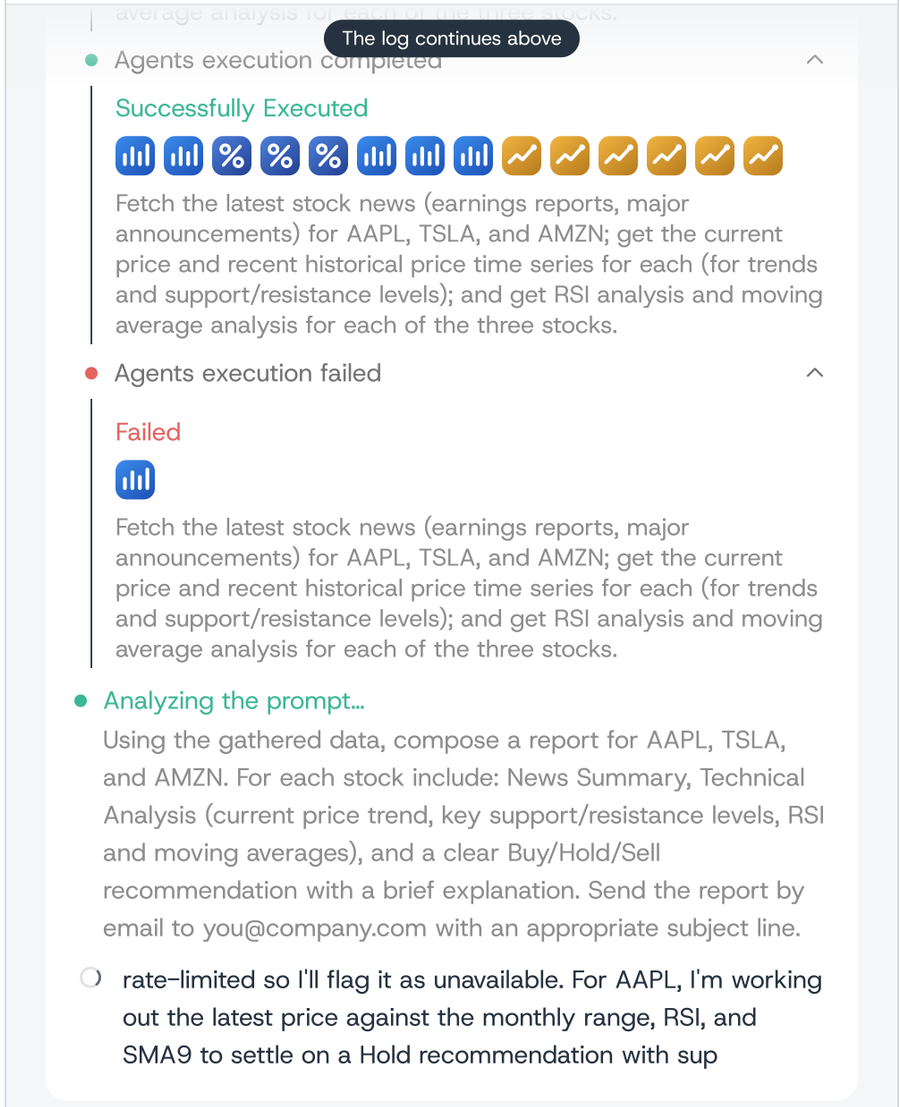
READ FROM THE STREAM

- `agents_retrieved` at 18.1 s with `retrievedAgents[]`
- `executing` at 18.1 s

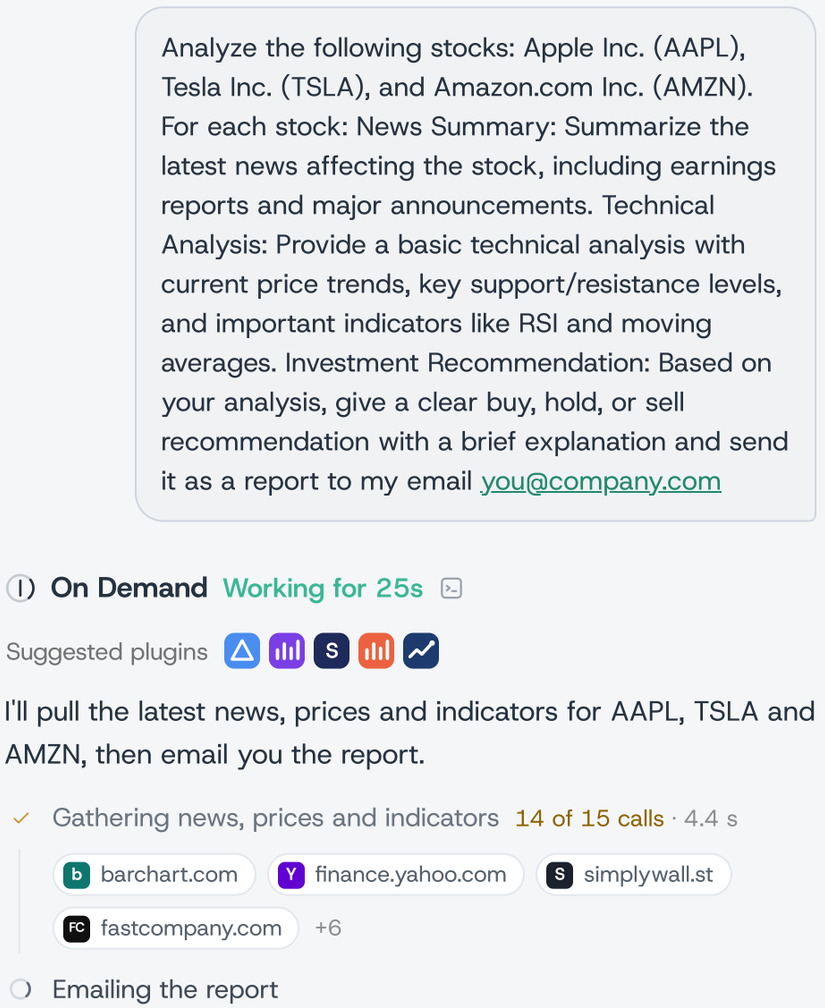
Results arrive as favicons, and the failed call stays visible

Fourteen calls returned and one hit a rate limit. Today that is two more groups with the instruction. Proposed: “14 of 15 calls”, the sites that came back, and Sources filling on the right.

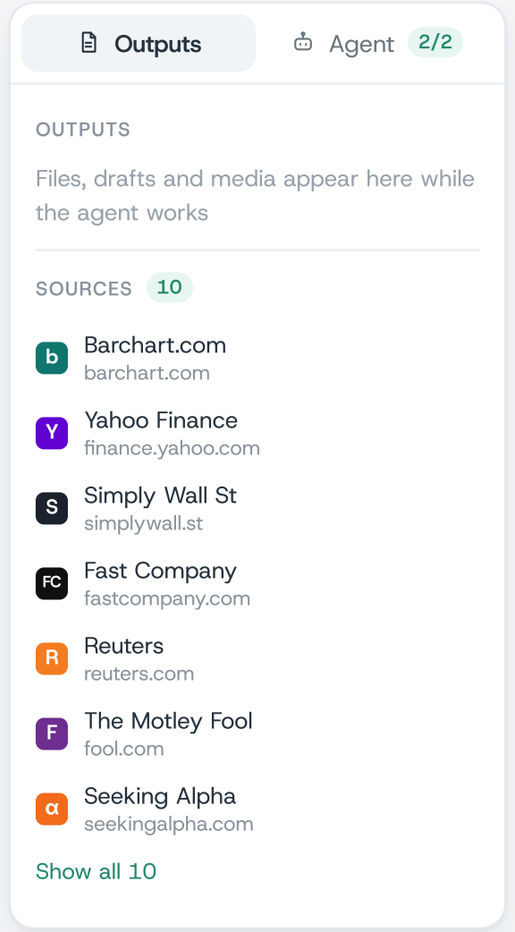
TODAY Every status frame becomes a row



PROPOSED The same run, with steps 1 and 2



PROPOSED Outputs card



TODAY DRAWS

- Group “Successfully Executed”: 14 logos and the instruction
- Group “Failed”: one logo and the instruction
- The next step starts with another full instruction

PROPOSED DRAWS

- The step shrinks to one line: “14 of 15 calls · 4.4 s” in amber
- Favicons of the returned pages appear one by one
- The Agent panel holds the 429 and its reason

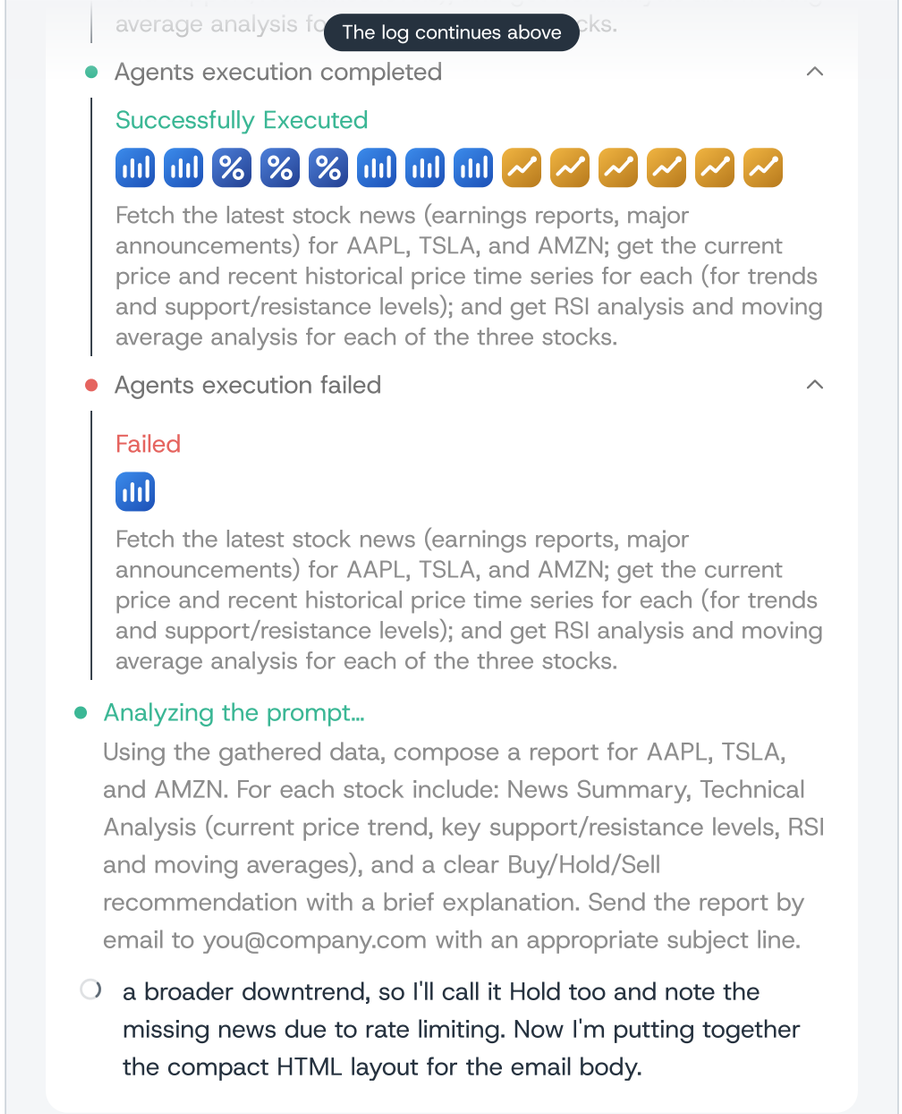
READ FROM THE STREAM

- `execution_completed` at 22.5 s: 14 `executedAgents[]` with `response`
- `execution_failed` at 22.5 s: `statusCode` 429, `errorResponse.message`

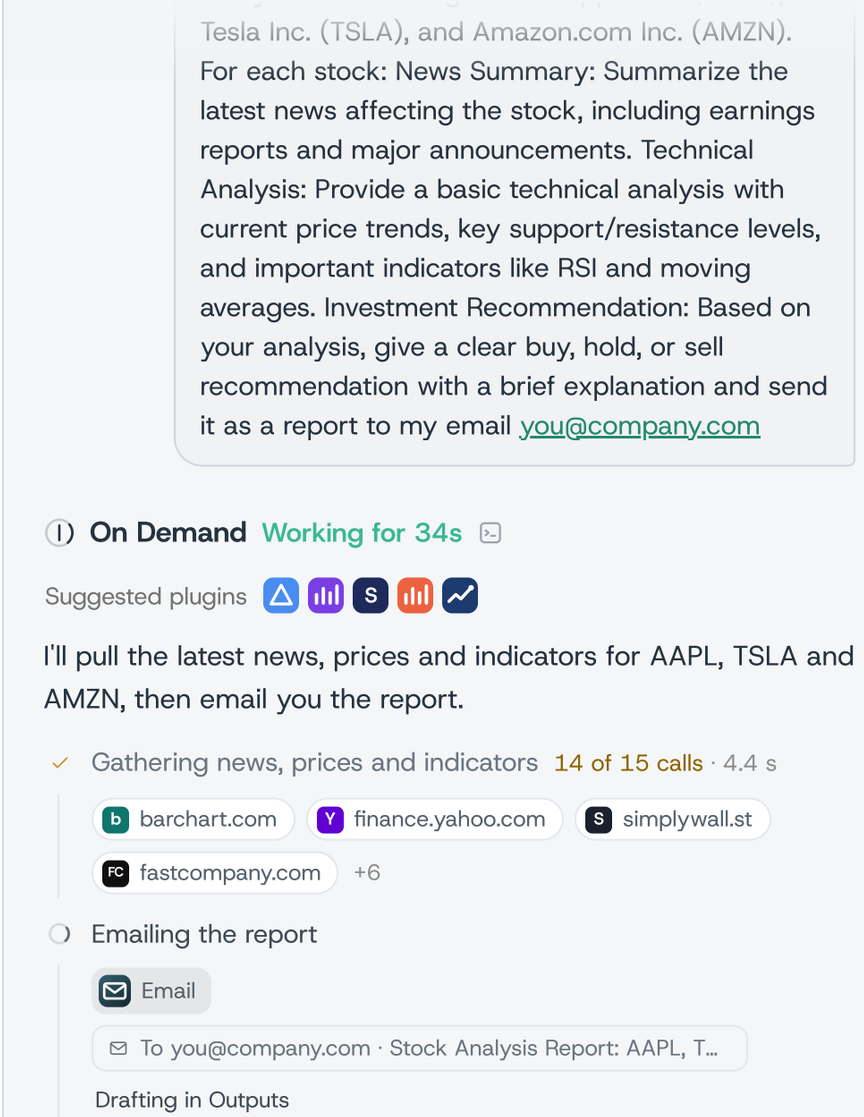
The email draft fills the Outputs card while it is being written

The model spends 13.4 s writing one call: the email body. Today nothing on screen changes during that time. Proposed: the step says what it is doing and the draft streams into Outputs.

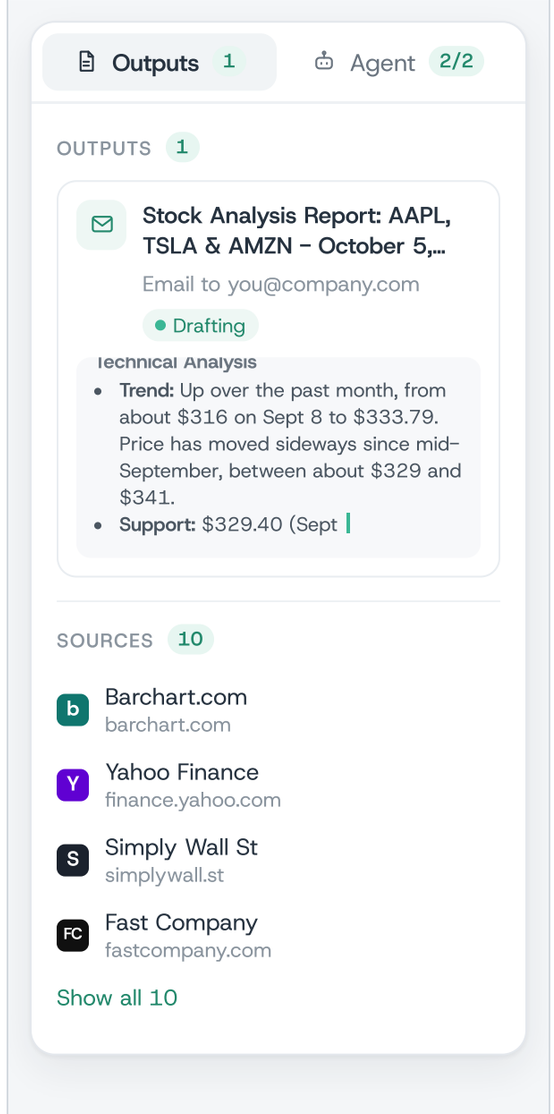
TODAY Every status frame becomes a row



PROPOSED The same run, with steps 1 and 2



PROPOSED Outputs card



TODAY DRAWS

- The second instruction, about 50 words
- Raw step reasoning
- The Outputs panel stays empty

PROPOSED DRAWS

- Step line “Emailing the report” with the recipient and the subject
- “Drafting in Outputs” points to the right
- The Outputs card shows the draft with a live preview

READ FROM THE STREAM

- `step_output` 27.9 to 41.3 s
- `api_request_parameters` : `to` , `subject` , then `text` as it streams

Twenty-two seconds of reasoning become three words

Between “fulfilling” and the first answer token the model reasons for 22 s. Today that is a second Thinking card with raw text. Proposed: “Writing the answer”, with the reasoning one click away.

TODAY Every status frame becomes a row

email to you@com… The log continues above …opriate subject line.

- Executing the agents…

Using the gathered data, compose a report for AAPL, TSLA, and AMZN. For each stock include: News Summary, Technical Analysis (current price trend, key support/resistance levels, RSI and moving averages), and a clear Buy/Hold/Sell recommendation with a brief explanation. Send the report by email to you@company.com with an appropriate subject line.

- Agents execution completed

Successfully Executed

Using the gathered data, compose a report for AAPL, TSLA, and AMZN. For each stock include: News Summary, Technical Analysis (current price trend, key support/resistance levels, RSI and moving averages), and a clear Buy/Hold/Sell recommendation with a brief explanation. Send the report by email to you@company.com with an appropriate subject line.

- Execution log created
- Fulfilling the prompt…

Thinking

couldn't be pulled due to a rate limit, so any AMZN commentary relies on general knowledge rather than fresh sources.
Now I'm pulling together the price data: AAPL, TSLA, and AMZN closing prices and recent movements to work through the technical picture fo

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 52s

Suggested plugins

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

- ✓ Gathering news, prices and indicators 14 of 15 calls · 4.4 s
  - barchart.com · finance.yahoo.com · simplywall.st · fastcompany.com +6
- ✓ Emailing the report Sent to you@company.com · 0.4 s
- Writing the answer ›

PROPOSED Outputs card

Outputs 1 | Agent 2/2

OUTPUTS 1

Stock Analysis Report: AAPL, TSLA & AMZN – October 5,…
Email to you@company.com
✓ Sent

Stock Analysis Report: AAPL, TSLA & AMZN
Report date: October 5, 2026 (data as of the October 2, 2026 close)
1. Apple Inc. (AAPL) – Price: $333.79

SOURCES 10

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com

Show all 10

TODAY DRAWS

- Rows: “Execution log created” and “Fulfilling the prompt…”
- A second Thinking card streaming raw reasoning and memory status lines

PROPOSED DRAWS

- Live line “Writing the answer”. A chevron opens the reasoning
- The email step reads “Sent to” and the address
- Outputs marks the draft Sent

READ FROM THE STREAM

- `execution_completed` at 41.8 s for the email
- `fulfilling` at 42.4 s
- `fulfillment_thinking` 42.4 to 64.4 s

When the first answer token lands, the steps fold into the header

Today the answer starts under the whole log. Proposed: the steps collapse into the header line, a chevron reopens them, and each link in the answer carries its favicon.

TODAY Every status frame becomes a row

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand

Thinking

- Finding relevant plugins...
- Plugin suggestions completed
  - Alchemy Agent
  - Chat Analytics Tool
  - Sentiment Analysis Report Tool
  - Quarterly financial report data tool
  - US Mutual Fund Technical
- Initializing the process...
- Analyzing the prompt...
- Execution plan created

The log continues. The answer starts further down

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 1m 10s

Stock Analysis Report: AAPL, TSLA, AMZN

The report has been sent to **you@company.com**, and the email service confirmed delivery. Prices below are as of the October 2, 2026 close. Support and resistance levels are drawn from recent closing prices over the past month.

Data gap: The Amazon news request failed because the data provider's rate limit was exceeded, so I have no current AMZN headlines. The AMZN section is based on technical data only, and I haven't filled the gap with unverified news.

1. Apple Inc. (AAPL): Recommendation: HOLD

PROPOSED Outputs card

Outputs 1 | Agent 2/2

OUTPUTS 1

Stock Analysis Report: AAPL, TSLA & AMZN – October 5,...
Email to you@company.com
Sent

Stock Analysis Report: AAPL, TSLA & AMZN
Report date: October 5, 2026 (data as of the October 2, 2026 close)
1. Apple Inc. (AAPL) – Price: $333.79

SOURCES 12

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com
- Reuters reuters.com · cited
- The Motley Fool fool.com
- Seeking Alpha

TODAY DRAWS

- The log stays open above the answer
- In the live app the answer begins 1,640 px below the header

PROPOSED DRAWS

- The steps fold within 200 ms of the first chunk
- The steps and “Thought for 33s” sit behind the chevron
- Answer links show a favicon and join Sources as cited

READ FROM THE STREAM

- `fulfillment` 64.4 to 83.3 s: `answer` chunks
- Markdown links inside the answer

At the end, the header and the answer are all that remain

Today 17 rows and 2 Thinking cards stay above the answer after the run. Proposed: the header keeps time and tokens, one chevron reopens the steps, and a sources button sits under the answer.

TODAY Every status frame becomes a row

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand 1m 20s 25662 tokens

Thinking

- Finding relevant plugins...
- Plugin suggestions completed
  - Alchemy Agent
  - Chat Analytics Tool
  - Sentiment Analysis Report Tool
  - Quarterly financial report data tool
  - US Mutual Fund Technical
- Initializing the process...
- Analyzing

The log continues. The answer starts further down

PROPOSED The same run, with steps 1 and 2

Analyze the following stocks: Apple Inc. (AAPL), Tesla Inc. (TSLA), and Amazon.com Inc. (AMZN). For each stock: News Summary: Summarize the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand 1m 20s · 25,662 tokens

Stock Analysis Report: AAPL, TSLA, AMZN

The report has been sent to **you@company.com**, and the email service confirmed delivery. Prices below are as of the October 2, 2026 close. Support and resistance levels are drawn from recent closing prices over the past month.

Data gap: The Amazon news request failed because the data provider's rate limit was exceeded, so I have no current AMZN headlines. The AMZN section is based on technical data only, and I haven't filled the gap with unverified news.

PROPOSED Outputs card

Outputs 1 | Agent

OUTPUTS 1

Stock Analysis Report: AAPL, TSLA & AMZN – October 5,...
Email to you@company.com
Sent

Stock Analysis Report: AAPL, TSLA & AMZN
Report date: October 5, 2026 (data as of the October 2, 2026 close)
1. Apple Inc. (AAPL) – Price: $333.79

SOURCES 16

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com
- Reuters reuters.com · cited
- The Motley Fool fool.com · cited

TODAY DRAWS

- 17 status rows and 2 Thinking cards stay above the answer
- Measured in the live app: 1,640 px of log on a 962 px screen

PROPOSED DRAWS

- Header: “1m 20s · 25,662 tokens”
- One chevron reopens the steps
- Outputs keeps the sent email and every source

READ FROM THE STREAM

- `fulfillment_completed` at 83.4 s with the full `answer`
- `metricsLog` at 83.4 s
- `[DONE]` at 85.3 s

Step 1 names the plugins in use. The prompt addenda name the task

With the frontend change alone, the opening sentence is today's plan objective and step labels name the plugins in use. With the planner addendum, the sentence turns first person and each step carries a label written for the user.

STEP 1 UI only · stock run, 0:34

including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 34s

Suggested plugins

Analyze AAPL, TSLA, and AMZN with a news summary, basic technical analysis (price trends, support/resistance, RSI, moving averages), and a buy/hold/sell recommendation for each, then email the report to you@company.com

✓ Using 3 plugins 14 of 15 calls · 4.4 s

barchart.com · finance.yahoo.com · simplywall.st · fastcompany.com · +6

Sending an email

Email

To you@company.com · Stock Analysis Report:...

Drafting in Outputs

STEP 2 With the addenda · 0:34

the latest news affecting the stock, including earnings reports and major announcements. Technical Analysis: Provide a basic technical analysis with current price trends, key support/resistance levels, and important indicators like RSI and moving averages. Investment Recommendation: Based on your analysis, give a clear buy, hold, or sell recommendation with a brief explanation and send it as a report to my email you@company.com

On Demand Working for 34s

Suggested plugins

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

✓ Gathering news, prices and indicators

14 of 15 calls · 4.4 s

barchart.com · finance.yahoo.com · simplywall.st · fastcompany.com · +6

Emailing the report

Email

To you@company.com · Stock Analysis Report:...

Drafting in Outputs

STEP 1 UI only · web search, 0:15

Fetch me images along with latest news about UAE from Perplexity tool

On Demand Working for 15s

Suggested plugins

Fetch images along with the latest news about UAE using the Perplexity tool

Searching the web

Perplexity

STEP 2 With the addenda · 0:15

Fetch me images along with latest news about UAE from Perplexity tool

On Demand Working for 15s

Suggested plugins

I'll search the web for the latest UAE news and pull an image for each story.

Searching the web for UAE news

Perplexity

STEP 1 LABELS

- "Using 3 plugins", "Sending an email"
- Built in the client from the plugins a step uses

STEP 2 LABELS

- "Gathering news, prices and indicators", "Emailing the report"
- Written by the planner as a `[label]` at the start of each step

STEP 1 SENTENCE

- The planner's `objective` as it is today, imperative and long

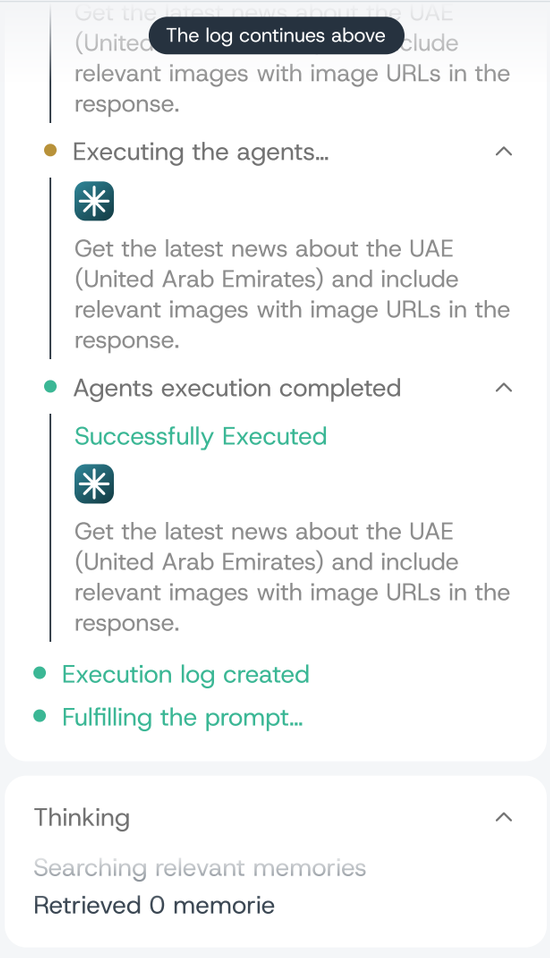
STEP 2 SENTENCE

- One first-person sentence, 22 words at most
- If a label is missing, the client falls back to the step 1 label

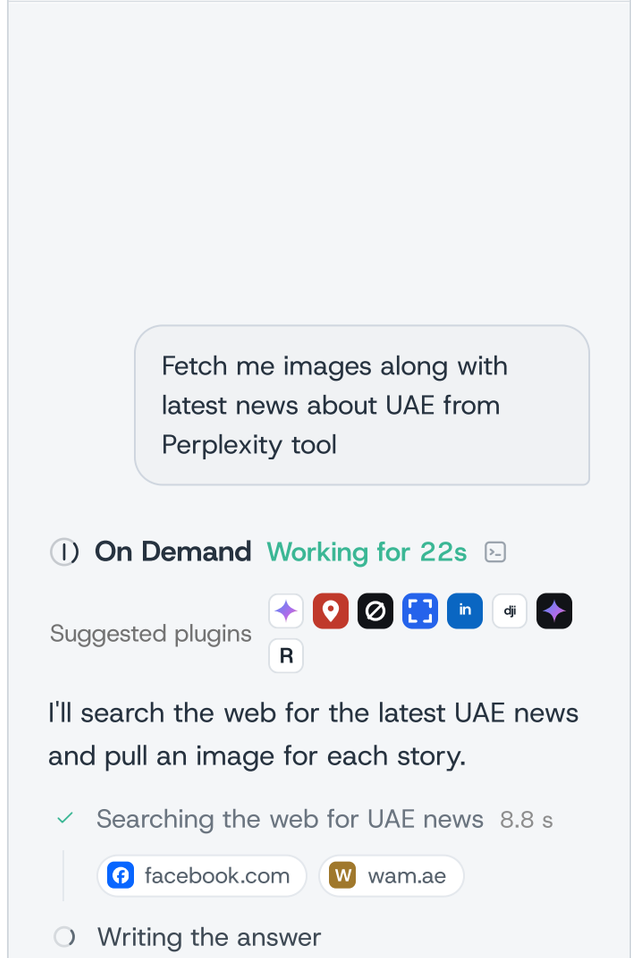
Search results show up as favicons: two when the search returns, the rest as the answer cites them

Today the search shows as three groups and no site at all. Proposed: one line, the sites as favicon chips, and a Sources list on the right. The plugin's response is cut at about 3,000 characters inside the event, so 2 sites are readable at this moment and the rest arrive as the answer cites them.

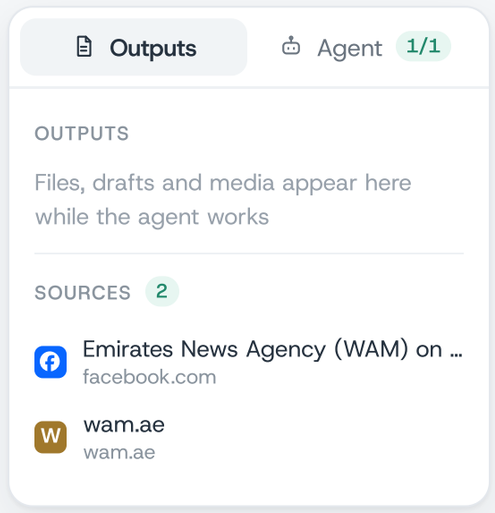
TODAY Three groups, no sites



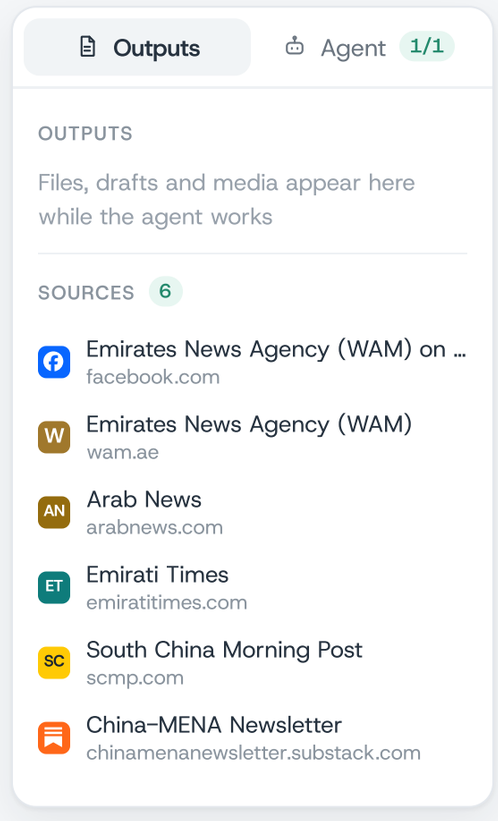
PROPOSED One line and the sites



PROPOSED Outputs card



STEP 3 Sources-first wrapper



TODAY DRAWS

- "Retrieved the agents", "Executing the agents..." and "Successfully Executed", each with the logo and the instruction
- No site, no link, no favicon until the answer

PROPOSED DRAWS

- "Searching the web for UAE news" with the Perplexity logo while it runs
- A favicon chip per page readable in the response, revealed 110 ms apart
- The search query itself stays out of the feed

READ FROM THE STREAM

- `execution_completed` at 19.6 s
- `executedAgents[].response`, scanned for links

OPTIONAL STEP 3

- The wrapper lists title and link first
- All sites arrive with the search result

The favicon trail through one search: searching, first sites, then every page the answer cites

Four moments of the proposed view in the same run. Favicons come from two places that already exist: links inside the plugin response, and markdown links in the answer as it streams.

1 0:15 · searching

Fetch me images along with latest news about UAE from Perplexity tool

On Demand Working for 15s

Suggested plugins

I'll search the web for the latest UAE news and pull an image for each story.

Searching the web for UAE news

Perplexity

SEARCHING

- The plugin chip pulses
- Three placeholder tiles hold the place of the sites to come

2 0:22 · first sites

Fetch me images along with latest news about UAE from Perplexity tool

On Demand Working for 22s

Suggested plugins

I'll search the web for the latest UAE news and pull an image for each story.

Searching the web for UAE news 8.8 s

facebook.com wam.ae

Writing the answer

FIRST SITES

- One favicon chip per page found in the response
- The step shows its duration: 8.8 s

3 0:35 · the answer cites more

Fetch me images along with latest news about UAE from Perplexity tool

On Demand Working for 35s

Latest UAE news with images

Placeholder answer text. The answer area is not changed by this proposal.

1. Story from the Emirates News Agency

Two-line summary of the story as returned by the search plugin, in the model's usual answer style. Read on WAM

2. Story from the WAM Facebook page

Two-line summary of the story as returned by the search plugin. Read the post

3. Story from Arab News

ANSWER LINKS

- The steps have folded into the header
- Each link in the answer carries its favicon and joins Sources

4 0:40 · done, steps reopened

Fetch me images along with latest news about UAE from Perplexity tool

On Demand 36s · 9,867 tokens

Suggested plugins

I'll search the web for the latest UAE news and pull an image for each story.

Searching the web for UAE news 8.8 s

facebook.com wam.ae

Thought for 15s

Latest UAE news with images

Placeholder answer text. The answer area is not changed by this proposal.

1. Story from the Emirates News Agency

Two-line summary of the story as returned by the search plugin, in the model's usual answer style. Read on WAM

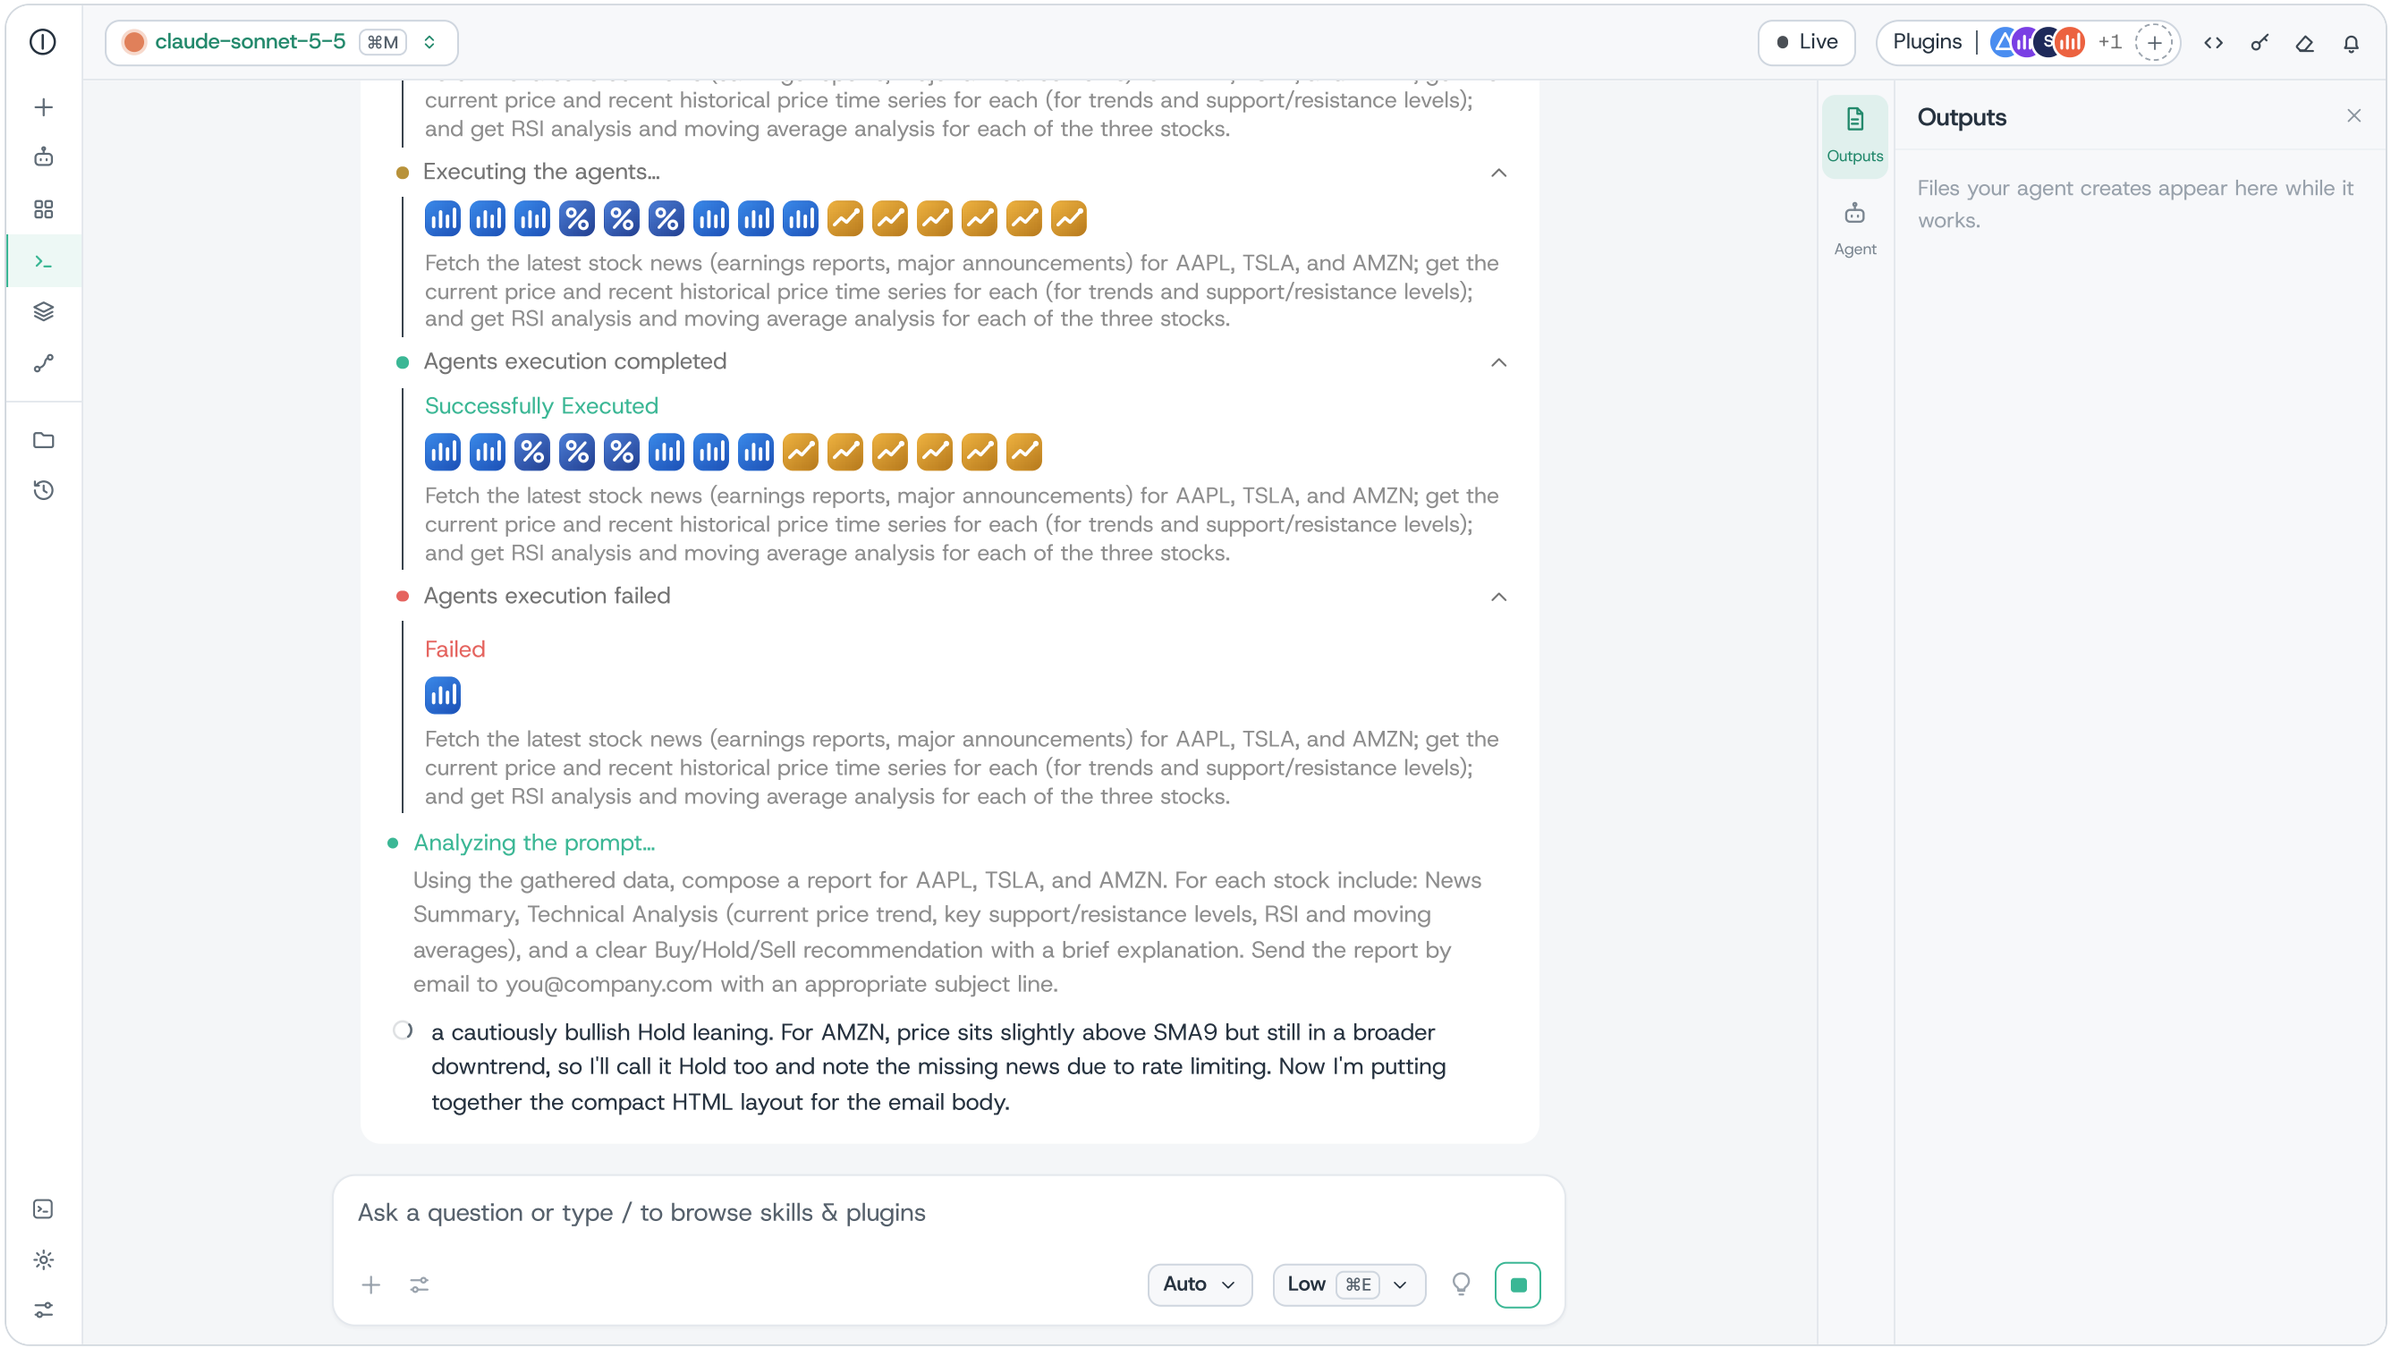
2. Story from the WAM Facebook page

DONE

- The chevron reopens the steps, with “Thought for” in place of the reasoning
- A sources button under the answer stacks the favicons

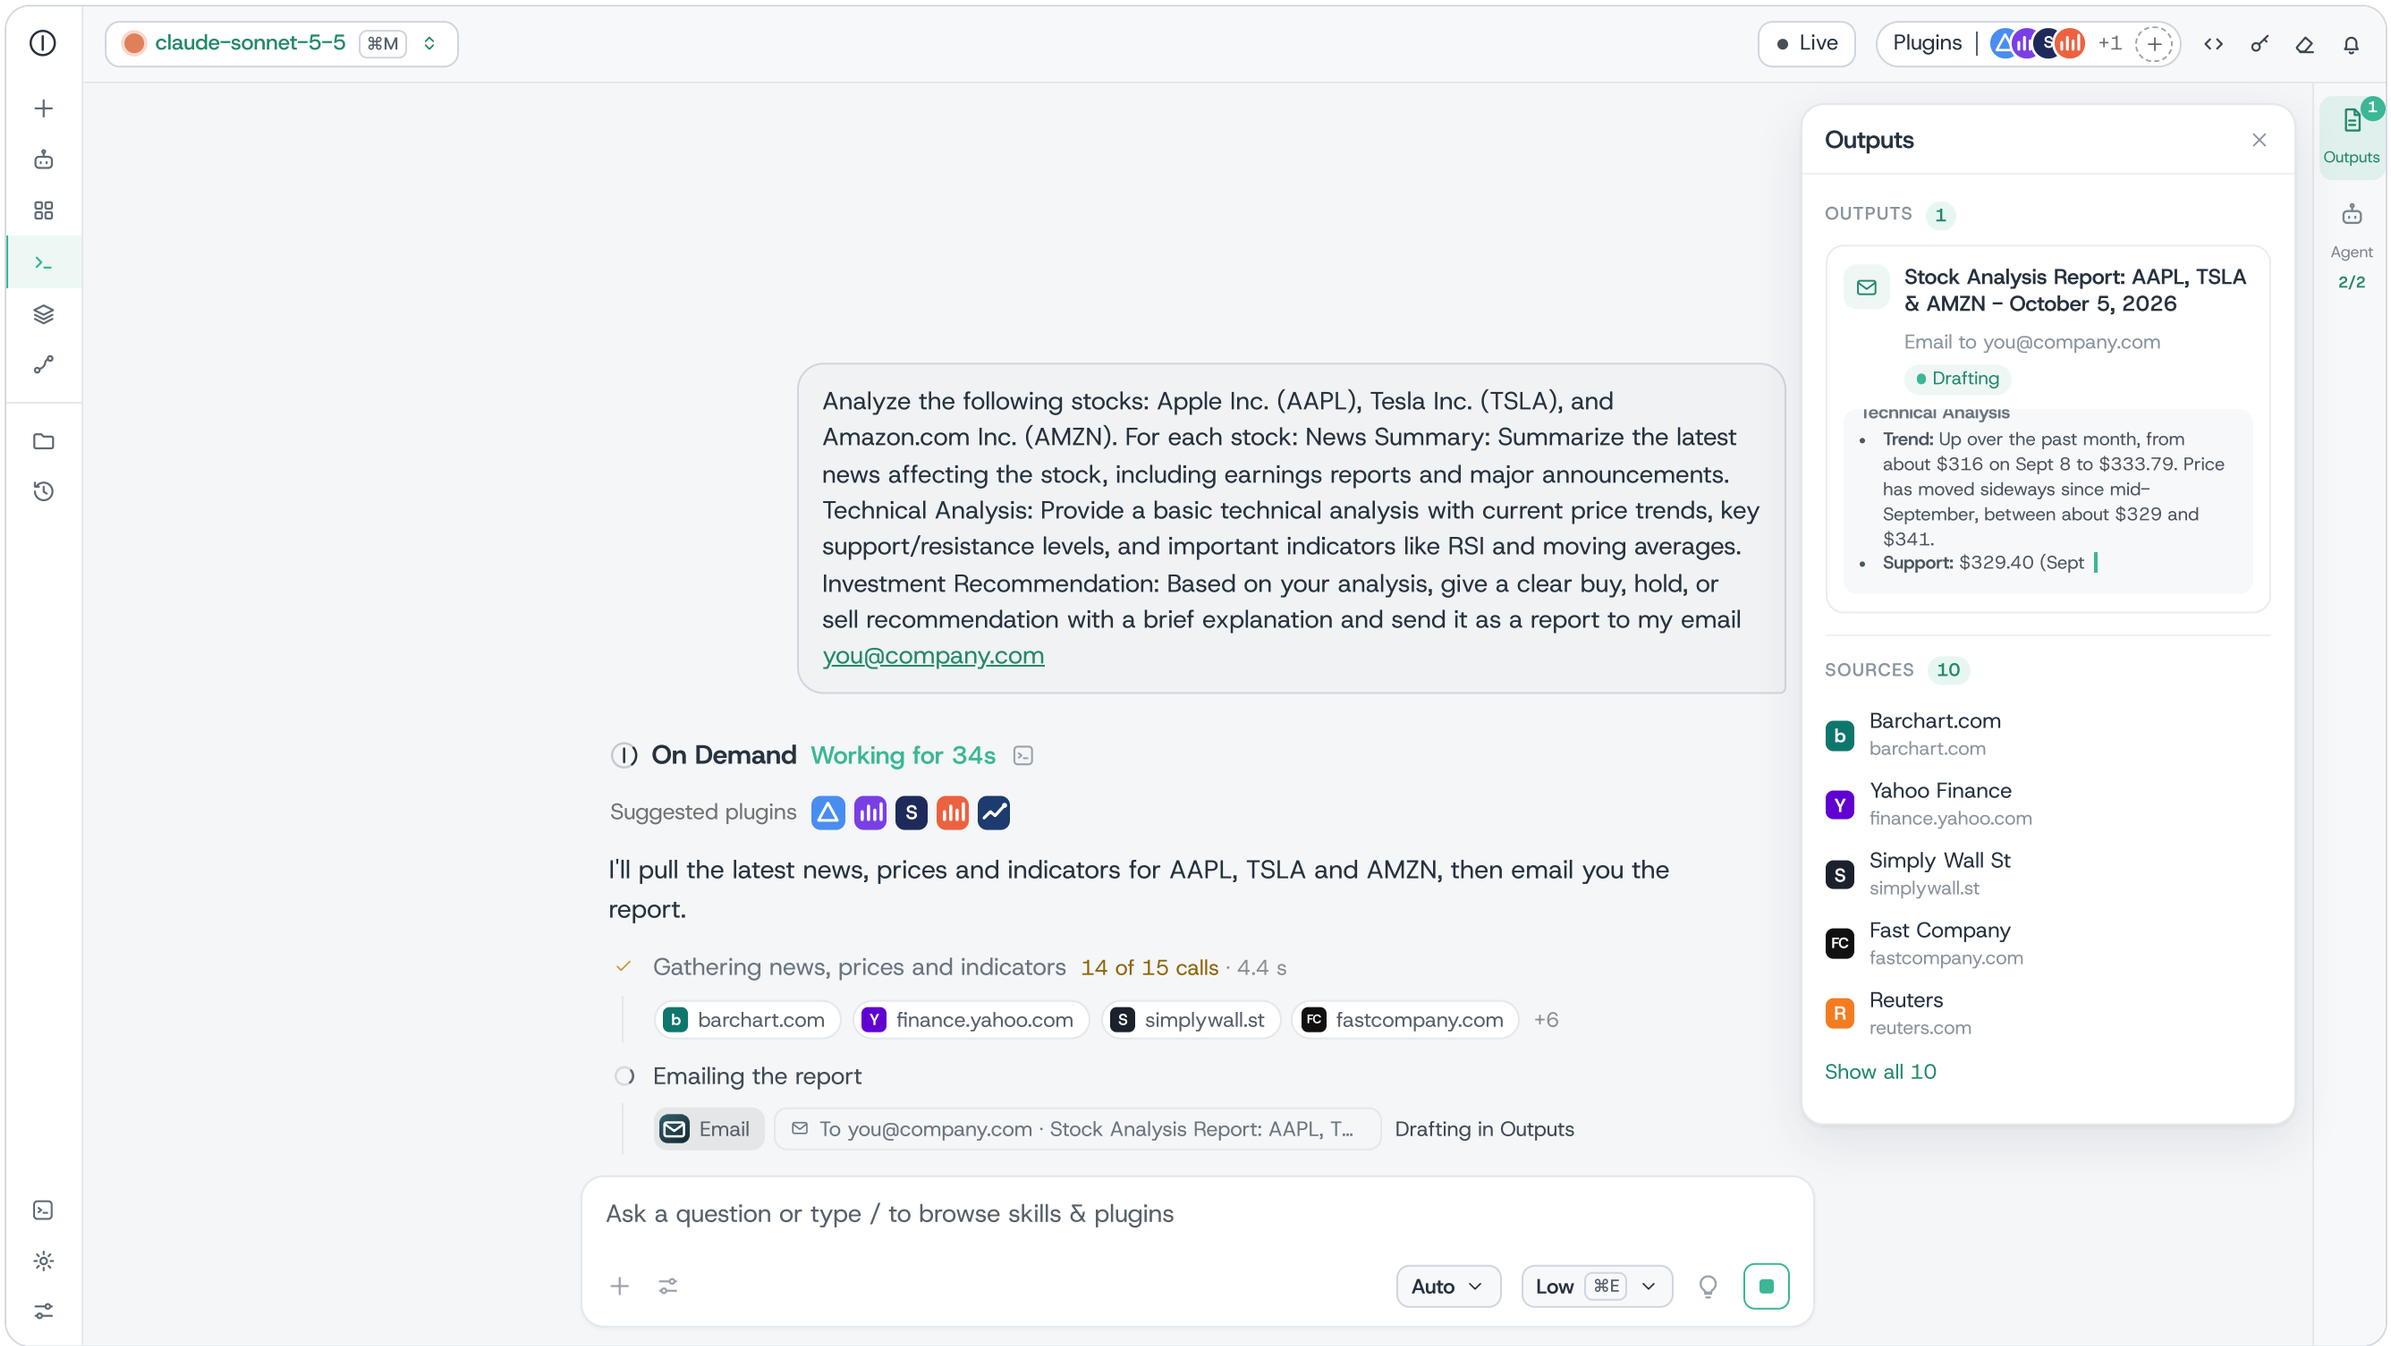
Today at 0:34: instructions and raw reasoning fill the column, and the right side is empty

The model is writing the email. Nothing on screen says so, and the Outputs panel has nothing to show.



Proposed at 0:34: six short lines in the column, the draft and the sources on the right

Same top bar, same composer, same rail buttons. The Outputs card floats over the gutter, so the conversation does not reflow. The Agent button shows step 2 of 2.



The right side fills while the agent works: outputs, sources, the plan, plugins and memory

Outputs lists what the run made and the pages it read. The Agent panel gains a “This run” section above the configuration that is there today. Both are fed by the same frames as the feed.

TODAY Outputs, whole run

Outputs

Files your agent creates appear here while it works.

0:25 Sources arrive

Outputs | Agent 2/2

OUTPUTS

Files, drafts and media appear here while the agent works

SOURCES 10

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com
- Reuters reuters.com
- The Motley Fool fool.com

Show all 10

0:34 The draft streams

Outputs 1 | Agent 2/2

OUTPUTS 1

Stock Analysis Report: AAPL, TSLA & AMZN – October 5,...
Email to you@company.com
Drafting

Technical Analysis
- **Trend:** Up over the past month, from about $316 on Sept 8 to $333.79. Price has moved sideways since mid-September, between about $329 and $341.
- **Support:** $329.40 (Sept

SOURCES 10

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com

Show all 10

1:25 Sent, 16 pages

Outputs 1 | Agent

OUTPUTS 1

Stock Analysis Report: AAPL, TSLA & AMZN – October 5,...
Email to you@company.com
Sent

Stock Analysis Report: AAPL, TSLA & AMZN
Report date: October 5, 2026 (data as of the October 2, 2026 close)

1. Apple Inc. (AAPL) – Price: $333.79

SOURCES 16

- Barchart.com barchart.com
- Yahoo Finance finance.yahoo.com
- Simply Wall St simplywall.st
- Fast Company fastcompany.com

Show all 16

0:25 Agent panel

Outputs | Agent 2/2

THIS RUN

I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN, then email you the report.

- ✓ Gathering news, prices and indicators
- Emailing the report

PLUGINS 3

Stock Market News
6 calls · 1 failed · 3.2 s
`getStockNews AMZN` 429 · rate limit

US Stock Fundamental Analysis
3 calls · 0.8 s ✓

US Stock Technical Analysis
6 calls · 0.8 s ✓

MEMORY

0 memories used
Checked 3 times

CONFIGURATION

- AGENT
- MODEL
- LIVE
- SAMPLING
- SKILLS
- ADVANCED

TODAY
- Empty in plugin runs
- Files show only in agent runs

SOURCES
- A row per page: favicon, title, site
- “cited” once the answer links it

DRAFT
- Recipient and subject first
- The body follows the stream

DONE
- Drafting turns to Sent
- A failure turns it to Failed

AGENT
- The sentence, the steps, plugins with call counts
- The failed call with 429 and its reason

Every frame keeps its name and its fields. Only what the client draws for it changes

SSE FRAME	TODAY DRAWS	PROPOSED DRAWS
`message · statusLog · plugin_suggestion.initialized`	Row: “Finding relevant plugins…”	Live line: Finding plugins
`message · statusLog · plugin_suggestion.completed`	Row, plus one chip per plugin with logo and name	“Suggested plugins” and logos only
`message · statusLog · initializing`	Row: “Initializing the process…”	No new line. The live line reads Thinking
`message · statusLog · analyzing`	Row: “Analyzing the prompt…”	No new line
`thinking · planning_thinking`	“Thinking” card with raw reasoning	Kept behind Thinking. Memory lines go to the Agent panel
`thinking · planning_output`	Not drawn	Read while it streams: the objective becomes the opening sentence
`message · statusLog · plan_created`	Row: “Execution plan created”	Confirms the sentence. Steps listed in the Agent panel
`message · statusLog · analyzing + stepQuery`	Row, plus the full step instruction	New step line with a short label
`thinking · step_thinking`	Raw reasoning under the row	Hidden. The step label shimmers
`thinking · step_output`	Not drawn	Each plugin call appears as a logo chip. Long text fills Outputs
`message · statusLog · agents_retrieved`	Group: logos and the instruction again	No new line. Chips become the confirmed call list
`message · statusLog · executing`	Group: logos and the instruction again	Chips pulse. The line reads “15 calls running”
`message · statusLog · execution_completed`	Group “Successfully Executed”: logos and the instruction again	Chips tick. Favicons of returned pages appear one by one
`message · statusLog · execution_failed`	Group “Failed”: logo and the instruction again	The step reads “14 of 15 calls”. Detail in the Agent panel
`message · statusLog · execution_log_created`	Row: “Execution log created”	Not drawn. Kept for the Logs button
`message · statusLog · fulfilling`	Row: “Fulfilling the prompt…”	Live line: Writing the answer
`thinking · fulfillment_thinking`	Second “Thinking” card with raw reasoning	Kept behind Writing the answer
`message · fulfillment`	Answer text under the log	Answer text. The steps fold into the header. Each link adds a favicon
`message · statusLog · fulfillment_completed`	Row: “Fulfillment completed”	Not drawn. The full answer replaces the streamed text
`message · metricsLog`	Header: time and tokens	Header: time and tokens
`heartbeat`	Nothing	Nothing. The timer runs in the browser
`message · [DONE]`	Stop turns back into Send	Stop turns back into Send

Frame types named in the SSE reference but absent from the recordings (history summary, attachments, clarification, errors) are mapped in the technical spec and still need a test fixture. When the first `agent` frame arrives, the feed stops its live line and today's terminal and to-do view takes over, unchanged.

Three text-only prompt addenda upgrade the wording. No key, schema or event changes

Each block is appended to an existing prompt. The paste-ready text, the before and after JSON and the eval plan are in the technical spec.

1 · Planner

Writes `planning_output`, forwarded as `plan_created.queryPlan`

- `objective` becomes one first-person sentence starting with "I'll", 22 words at most
- Each `user_query` starts with a `[label]` of 3 to 7 words, then the instruction unchanged
- Nothing else changes: same keys, same order, same JSON

TODAY

```
"objective": "Analyze AAPL, TSLA, and AMZN with a news
summary, basic technical analysis (price trends, …),
and a buy/hold/sell recommendation for each, then
email the report to …"
"user_query": "Fetch the latest stock news (earnings
reports, …"
```

WITH THE ADDENDUM

```
"objective": "I'll pull the latest news, prices and
indicators for AAPL, TSLA and AMZN, then email you the
report."
"user_query": "[Gathering news, prices and indicators]
Fetch the latest stock news (earnings reports, …"
```

2 · Step executor

Writes `step_output`, the plugin calls JSON

- A fixed key order per call, so a chip can be drawn before the parameters finish
- Short parameters first (`to` , `subject` , `symbol`), long text last, so Outputs can title a draft at once
- Never copy the `[label]` into a parameter, subject line or message body

ALREADY TRUE IN THE RECORDING, NOW A RULE

```
{"pluginId": "plugin-1722285968", "name": "sendEmail",
…
 "api_request_parameters": {"to": "…", "subject":
"Stock Analysis Report: …", "text": "<the long body,
last>"}, …}
```

WHY IT MATTERS

The third rule guards the main risk of step 2: a label leaking into a sent email or a search query. The eval gate for that is zero.

3 · Answer

Writes `fulfillment`

- Cite each external fact with a markdown link to the exact page the plugin returned. Sources and favicons are built from these links
- Start with the result. The steps are already on screen
- One sentence starting with "Data gap:" when a call failed

CLOSE TO THIS IN THE RECORDING, NOW A RULE

```
… the stock jumped about 5% ([CNBC]
(https://www.cnbc.com/2026/10/02/…), [Reuters](…))

**Data gap:** The Amazon news request failed because
the data provider's rate limit was exceeded …
```

RECORDED

21 links on 16 sites in the stock answer, 9 links on 7 sites in the search answer.

GATE BEFORE STEP 2 SHIPS	VALUE
Eval set	30 prompts: 10 plugin runs, 10 web searches, 10 multi-step runs with a side effect
Valid plan and call JSON	100%, with the parser unchanged
Labels matching the pattern	95% or more
Labels copied into a plugin parameter	0
Answer quality	Equal or better in a blind pair comparison

WHAT THE PROMPTS SHOULD NOT DO

- Do not ask the model to narrate progress in its reasoning. Reasoning is hidden by default and its wording changes with the model
- Do not rename keys, and add none unless the plan parser is known to ignore them
- `objective` and `stepQuery` also reach API customers. Decide whether the addenda apply to all traffic or to the Playground first

The three blocks as they would be appended to the planner, step and answer prompts

Text only. If the planner forgets a label, the client falls back to a label built from the plugins in use, so nothing breaks.

1 · Planner addendum

Append to the planner system prompt

```
## Progress copy shown to the user

The interface shows two parts of your plan to the user while the work runs. Write them
for a person.

1. "objective" is ONE sentence in the first person, starting with "I'll", 22 words at
most.
   Say what you will do and what the user gets. Plain words. No plugin names, operation
ids, step numbers or JSON.
   Say "you" and "your" instead of repeating the user's name or email address.
   Good: "I'll pull the latest news, prices and indicators for AAPL, TSLA and AMZN,
then email you the report."
   Bad:  "Analyze AAPL, TSLA, and AMZN with a news summary, basic technical analysis
(price trends, support/resistance, RSI, moving averages), and a buy/hold/sell
recommendation for each, then email the report"

2. Every step's "user_query" starts with a progress label in square brackets, then one
space, then the full instruction exactly as you would otherwise write it.
   Label: 3 to 7 words, starts with a verb ending in -ing, names the thing being worked
on, no final punctuation, no plugin names or ids.
   Good: "[Gathering news, prices and indicators] Fetch the latest stock news (earnings
reports, major announcements) for AAPL, TSLA, and AMZN; ..."
   Good: "[Searching the web for UAE news] Get the latest news about the UAE (United
Arab Emirates) and ..."
   Good: "[Checking the facial recognition tool for cross-analysis] Run the uploaded
photo through the face match plugin and ..."
   Bad:  "[Step 1] ...", "[getStockNews x3] ...", "[Working] ..."

3. Change nothing else: same keys, same order, same JSON, same code fence.
```

2 · Step executor addendum

Append to the prompt that writes the plugin calls

```
## Output hygiene (the interface reads this JSON while it streams)

1. Keep this key order in every plugin object: pluginId, name, description,
api_request_parameters, all_parameters_hydrated, dependencies, identifier.
2. Inside api_request_parameters, write short identifying values first (to, recipient,
subject, title, symbol, query, message, url, filename) and long content last (text,
body, html, content, prompt).
3. A step query may start with a label in square brackets, such as "[Emailing the
report]". The label is display text for the user. Never copy it into any parameter,
subject line or message body.
```

3 · Answer addendum

Append to the answer prompt

```
## Sources and voice

1. Cite every external fact with an inline markdown link to the exact page URL the
plugin returned: ([Publisher](https://full.url/of/the/page)). Never cite a publisher by
name without its link when a link was returned. The interface builds the source list
and the favicons from these links.
2. Start with the result. Do not describe the steps you took. The interface already
shows them.
3. If a plugin call failed, say so once, in one sentence that starts with the bold
label "Data gap:". Then continue with what you have.
```

MONITOR, WITHOUT BLOCKING

- `objective` starts with “I'll” and has 22 words at most
- Each `user_query` matches `^\[[A-Z][a-z]+ing\b[^\]]{2,60}\] \S`
- No label text appears inside any `api_request_parameters` value

ROLLOUT AND ACCEPTANCE

Step 1 ships behind a flag that is also its kill switch. Steps 2 and 3 follow in either order

STEP 1 · UI ONLY · BEHIND THE FLAG `runViewV2`

0 · Reducer S
The stream reducer, two recorded fixtures and the unit tests. Nothing visible
Gate: tests green in CI

1a · Run feed M
The run header and the feed, staff only
Gate: feed checks pass on both fixtures

1b · Right rail M
The Outputs card and the Agent run section
Gate: output and source checks pass

1c · Ramp S
10% of Playground sessions, then 100%
Gate: measures hold for one week

THEN, IN EITHER ORDER

2 · Prompt addenda S
Planner and step executor together, then the answer block
Gate: the eval gates

3 · Sources first S
The search wrapper change
Gate: sources at execution equal the list length

Sizes are estimates for one frontend engineer who knows the Playground code. S is days, M is one to two weeks.

FLAG

- `runViewV2`, read once per session, off by default
- Off renders today's components untouched. Turning it off needs no deploy
- The Logs button opens the raw log in both states
- After step 2, the old view would print the `[label]` prefix when the flag is off. Strip it there with the same pattern

WHAT TO MEASURE

- Stop pressed before the first answer token: expected down
- Regenerate and thumbs down on plugin runs: flat or down
- Steps reopened, reasoning opened, sources clicked: shows what people still want to see
- After step 2, the share of steps without a planner label: under 5%

No baseline was available, so these are directions. Set targets after the first week at 10%.

ACCEPTANCE CHECKS · FIXTURES COVER THE STREAM CHECKS, THE LIVE APP COVERS THE REST

CHECK	PASSES WHEN	CHECK	PASSES WHEN
No server wording	No `statusMessage` is printed as it arrives	**Favicons**	First chip within 150 ms of `execution_completed`
No instruction text	`stepQuery` appears 0 times in the feed	**Failed call**	“14 of 15 calls”, with 429 in the Agent panel
Feed height	250 px or less and 6 lines or fewer at every moment	**Email output**	Draft while written, then “Sent to” the address
Suggested plugins	The label and logos only. Names on hover	**Fold**	Within 200 ms of the first answer chunk
Opening sentence	Starts typing before `plan_created`	**Agent runs**	Terminal, to-do list and files exactly as today
First plugin chip	Appears before `agents_retrieved`	**Reconnect**	Resume with `lastEventIndex` ends in the same state
Reasoning	Hidden unless opened. “Thought for 33s” after	**Flag off**	Today's components, untouched
Sources	The count equals the unique pages from responses plus answer links	**Access**	No shimmer under reduced motion. Every toggle works from the keyboard

Step 1 carries no backend risk. The real risks sit in the prompt changes

MAIN RISKS · EIGHT OF THE TEN IN THE SPEC

RISK	LIKELY	IMPACT	MITIGATION
A step label is copied into a plugin parameter, such as an email subject	Low	High	Executor rule 3. Eval gate of 0. Playground traffic first
Prompt changes reach API customers who print `objective` or `stepQuery`	Certain if prompts are shared	Medium	Decide the scope before step 2
The planner drops or garbles a label	Medium	Low	Plugin-based fallback label. Miss rate logged
`step_output` changes shape after a model update	Medium	Medium	Treated as optional. Chips then appear about 6 s later
Builders miss the full log while debugging	Medium	Medium	Logs button, reasoning one click away, calls and failures in the Agent panel
The 22 s before the answer still feels slow	Certain	Medium	“Writing the answer” and the timer. The fix is model latency
Most search sources stay hidden until the answer	Certain today	Low	Answer links fill the list. Step 3 removes the gap
Model-written HTML or a secret in a parameter is shown	Low	High	Sanitize previews. Mask secret keys. Never show plugin URLs

COUNTER-ARGUMENTS

THE CASE AGAINST	RESPONSE
The full log is our transparency	Nothing is removed. Plan, calls, failures and memory move to the Agent panel and the raw log stays behind Logs
A playground is for builders, who want raw frames	True for plugin debugging. The log, the reasoning and the status codes stay one click away, and a full-log switch is an open decision
Prompt changes risk regressions for cosmetic gain	Agreed, so step 1 ships without them. The fallback labels work if step 2 never ships
Favicons are decoration	They are the only sign of which sites were read before the answer arrives
A dedicated progress event would be cleaner	True in the long run. It was ruled out for now, and the streaming parser passed all 23,586 prefixes of the recorded documents
This copies ChatGPT	It adopts conventions users already know. The terminal view, plugin logos and the rail stay OnDemand's own

TRADE-OFFS ACCEPTED

Less on screen, one click for the rest
Six lines instead of the full log. Builders pay one click to see the detail

Client-side reading instead of a new event
The client reads JSON while it streams. Cheaper now, less clean than a dedicated event later

Labels name plugins until step 2
“Using 3 plugins” names the tools, not the task. The planner label fixes that at the price of a prompt change

Sources are partial until the answer
2 of 7 search sites when the search returns, unless step 3 ships

Six decisions are open, and none of them blocks step 1

OPEN DECISIONS

DECISION	OPTIONS	RECOMMENDED
Does the Outputs card open by itself?	Never, once per session on the first output, or every run	Once per session, then remember the user's choice
Where does today's full log live?	Logs button only, or also a "Show full log" switch under Advanced	Logs only. Add the switch if builders ask
Who gets the prompt addenda?	All traffic, Playground traffic only, or nobody	Playground first, then all traffic once copied labels stay at 0
How does the step label travel?	A bracket prefix in `user_query`, or a `label` key in the planner JSON	The prefix now. The key if the plan parser ignores unknown keys
Show the search query as a chip, as ChatGPT does?	Off, or on for search plugins	Off. The brief asked for no plugin query in the feed
Where do favicons load from?	Each site directly, a proxy, or off	Directly, with a letter-tile fallback

NEXT STEPS

ACTION	OWNER
Settle the six decisions	Mk
Walk through the replay and fix the copy	Design
Port the reducer, fixtures and tests into the Playground repo	Frontend
Record three more fixtures: an agent-mode run, a failed plan and a clarification	Frontend
Ask the RAG owner whether the plan parser ignores unknown keys and whether prompts can differ by route	Prompt owner
Run the 30-prompt eval on the three addenda	Prompt owner

ASSUMPTIONS

- Both recordings ran in standard chat mode at Low reasoning effort on claude-sonnet-5-5. Other models, effort levels and agent mode were not recorded
- The Playground sends `debugMode: "on"`. Whether the `thinking` channel depends on it was not tested
- The client resolves plugin names and logos from its own catalog, as it does today for retrieved agents
- The planner, step and answer prompts were not available. The addenda are written against their recorded output
- The cut at about 3,000 characters is inferred from responses ending in "… (truncated)"
- Scores, sizes and likelihoods are estimates

WHAT IS REAL AND WHAT IS ILLUSTRATIVE

- Event names, fields, order and timings come from the two recordings
- Proposed panes show the step 2 wording: the sentence and the step labels are written as the planner addendum would write them. Page 15 shows step 1 alone
- The panes run on a thinned copy of each recording: 726 of 1,033 frames and 203 of 414, with one search site left out
- Long texts are abridged, most article links are shortened to the publisher's home page and the email address is a placeholder
- The web-search answer body is placeholder text, since the answer area does not change
- Logos and favicons are drawn stand-ins. Production uses `logoUrl` and each site's own favicon
- The 1,640 px figure was measured in the live app. Proposed heights were measured on the replica at a 712 px column

SOURCES

- OnDemand Chat Streaming Reference and Submit Query v1 SSE event reference, supplied PDFs
- Two streams recorded on app.on-demand.io/playground, 5 Oct 2026: 1,033 and 414 frames
- 14 supplied screenshots. ChatGPT and Claude web apps, checked live on 5 Oct 2026
- Technical spec with the reducer, component specs and the full prompt text: OnDemand Playground Run View: Technical Spec